```
Last login: Sun Apr 12 11:17:28 on ttys000
Run-Mac:~ mac$ cd ~/.ssh
Run-Mac:.ssh mac$ ssh -i "Runzhe.pem" ubuntu@ec2-3-80-160-96.compute-1.amazonaws.com
ssh: connect to host ec2-3-80-160-96.compute-1.amazonaws.com port 22: Connection refused
Run-Mac:.ssh mac$ ssh -i "Runzhe.pem" ubuntu@ec2-3-80-160-96.compute-1.amazonaws.com
Warning: Permanently added the ED25519 host key for IP address '3.80.160.96' to the list of known hosts.
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1060-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

  System information as of Sun Apr 12 16:03:32 UTC 2020

  System load:  0.72               Processes:           226
  Usage of /:   28.0% of 30.96GB   Users logged in:     0
  Memory usage: 1%                 IP address for ens5: 172.31.0.69
  Swap usage:   0%

 * Kubernetes 1.18 GA is now available! See https://microk8s.io for docs or
   install it with:

     sudo snap install microk8s --channel=1.18 --classic

 * Multipass 1.1 adds proxy support for developers behind enterprise
   firewalls. Rapid prototyping for cloud operations just got easier.

     https://multipass.run/

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

89 packages can be updated.
39 updates are security updates.


Last login: Fri Apr  3 19:45:17 2020 from 107.13.161.147
export openblas_num_threads=1; export OMP_NUM_THREADS=1; python EC2.py
ubuntu@ip-172-31-0-69:~$ export openblas_num_threads=1; export OMP_NUM_THREADS=1; python EC2.py
12:05, 04/12; num of cores:16
large_A_vary_sd_R

Basic setting:[rep_times, sd_O, sd_D, sd_u_O, w_O, w_A, u_O_u_D, sd_R_range, t_func] = [16, None, None, 20, 0.5, 1.5, 0, [0, 10, 20, 30]
, None]


---------------------------------------
[pattern_seed, day, sd_R] = [2, 7, 0]

max(u_O) =  145.8
O_threshold = 100
number of reward locations:  9
O_threshold = 105
number of reward locations:  7
O_threshold = 110
number of reward locations:  6
O_threshold = 115
number of reward locations:  3
O_threshold = 125
number of reward locations:  2
target 1 in 5 DONE!
^CTraceback (most recent call last):
Process Process-4:
  File "EC2.py", line 81, in <module>
    with_MF = with_MF, with_NO_MARL = with_NO_MARL, with_IS = with_IS,
Process Process-5:
  File "/home/ubuntu/simu_funs.py", line 63, in simu
Process Process-14:
Process Process-12:
    value_reps = parmap(once, range(OPE_rep_times), n_cores)
  File "/home/ubuntu/_uti_basic.py", line 80, in parmap
    [q_in.put((None, None)) for _ in range(nprocs)]
  File "/home/ubuntu/_uti_basic.py", line 80, in <listcomp>
    [q_in.put((None, None)) for _ in range(nprocs)]
  File "/home/ubuntu/anaconda3/lib/python3.7/multiprocessing/queues.py", line 82, in put
Process Process-16:
Process Process-9:
Process Process-8:
Process Process-11:
Process Process-1:
    if not self._sem.acquire(block, timeout):
KeyboardInterrupt
ubuntu@ip-172-31-0-69:~$ export openblas_num_threads=1; export OMP_NUM_THREADS=1; python EC2.py
12:21, 04/12; num of cores:16
large_A_vary_sd_R

Basic setting:[rep_times, sd_O, sd_D, sd_u_O, w_O, w_A, u_O_u_D, sd_R_range, t_func] = [16, None, None, 20, 0.5, 1.5, 0, [0, 10, 20, 30]
, None]
```

```
--------------------------------------
[pattern_seed, day, sd_R] = [2, 7, 0]

max(u_O) =  145.8
O_threshold = 125
number of reward locations:  2
target 1 in 1 DONE!

--------------------------------------
Value of Behaviour policy:55.239
O_threshold = 125
MC for this TARGET:[68.265, 0.065]
   [DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]
bias:[[-13.91, -13.91, -13.96]][[-20.03, -68.26, -13.03]]
std:[[0.74, 0.73, 0.48]][[0.37, 0.0, 0.23]]
MSE:[[13.93, 13.93, 13.97]][[20.03, 68.26, 13.03]]
MSE(-DR):[[0.0, 0.0, 0.04]][[6.1, 54.33, -0.9]]
***
==============
******************** THIS SETTING IS GOOD ********************
[[13.93 13.93 13.97 20.03 68.26 13.03]]


time spent until now: 10.3 mins

12:31, 04/12

--------------------------------------
[pattern_seed, day, sd_R] = [2, 7, 10]

max(u_O) =  145.8
O_threshold = 125
number of reward locations:  2
```