

```
Last login: Sun Mar 29 15:26:19 on ttys000
Run-Mac:~ mac$ cd ~/.ssh
Run-Mac:~.ssh mac$ ssh -i "Runzhe.pem" ubuntu@ec2-35-171-129-20.compute-1.amazonaws.com
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1060-aws x86_64)
```

```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:        https://ubuntu.com/advantage
```

System information disabled due to load higher than 16.0

```
* Kubernetes 1.18 GA is now available! See https://microk8s.io for docs or
install it with:
```

```
sudo snap install microk8s --channel=1.18 --classic
```

```
* Multipass 1.1 adds proxy support for developers behind enterprise
firewalls. Rapid prototyping for cloud operations just got easier.
```

```
https://multipass.run/
```

```
* Canonical Livepatch is available for installation.
- Reduce system reboots and improve kernel security. Activate at:
https://ubuntu.com/livepatch
```

```
50 packages can be updated.
0 updates are security updates.
```

```
*** System restart required ***
Last login: Sun Mar 29 19:26:23 2020 from 107.13.161.147
ubuntu@ip-172-31-4-46:~$ export openblas_num_threads=1; export OMP_NUM_THREADS=1
ubuntu@ip-172-31-4-46:~$ python EC2.py
15:32, 03/29; num of cores:16
```

```
Basic setting:[sd_0, sd_D, sd_R, sd_u_0, w_0, w_A, lam] = [2, 2, None, 0.4, 1, 1, 0.0001]
```

```
-----
[pattern_seed, T, sd_R] = [0, 672, 0]
```

```
max(u_0) = 27.327727595549877
0_threshold = 12
means of Order:
```

```
22.323 12.937 16.305 27.014 23.267
```

```
7.457 16.12 10.376 10.577 12.991
```

```
11.677 19.721 14.946 11.573 13.165
```

```
12.597 20.038 10.155 12.494 7.833
```

```
3.97 14.317 15.577 8.192 27.328
```

```
target policy:
```

```
1 1 1 1 1
```

```
0 1 0 0 1
```

```
0 1 1 0 1
```

```
1 1 0 1 0
```

```
0 1 1 0 1
```

```
number of reward locations: 16
0_threshold = 9
target policy:
```

```
1 1 1 1 1
```

```
0 1 1 1 1
```

```
1 1 1 1 1
```

```
1 1 1 1 0
```

0 1 1 0 1

number of reward locations: 21

0_threshold = 15

target policy:

1 0 1 1 1

0 1 0 0 0

0 1 0 0 0

0 1 0 0 0

0 0 1 0 1

number of reward locations: 9

1 2 3 1 2 3

0_threshold = 12

MC-based mean and std of average reward:[1.1718e+01 5.0000e-03]

Value of Behaviour policy:11.24

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.11, 0.11, 0.1]][[0.15, 0.14, 0.15]][[11.72, 11.72, 11.72]][[0.1, 0.48]]
std:[[0.0, 0.0, 0.0]][[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][[0.0, 0.01]]
MSE:[[0.11, 0.11, 0.1]][[0.15, 0.14, 0.15]][[11.72, 11.72, 11.72]][[0.1, 0.48]]
MSE(-DR):[[0.0, 0.0, -0.01]][[0.04, 0.03, 0.04]][[11.61, 11.61, 11.61]][[-0.01, 0.37]]
better than DR_NO_MARL

0_threshold = 9

MC-based mean and std of average reward:[1.1523e+01 5.0000e-03]

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.4, 0.39, 0.4]][[0.45, 0.44, 0.45]][[11.52, 11.52, 11.52]][[0.39, 0.28]]
std:[[0.03, 0.03, 0.01]][[0.02, 0.02, 0.02]][[0.0, 0.0, 0.0]][[0.01, 0.01]]
MSE:[[0.4, 0.39, 0.4]][[0.45, 0.44, 0.45]][[11.52, 11.52, 11.52]][[0.39, 0.28]]
MSE(-DR):[[0.0, -0.01, 0.0]][[0.05, 0.04, 0.05]][[11.12, 11.12, 11.12]][[-0.01, -0.12]]
***** BETTER THAN [QV, IS, DR_NO_MARL] *****

MC-based ATE = -0.2

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]
bias:[[0.29, 0.28, 0.3]][[0.3, 0.3, 0.3]][[0.2, 0.2, 0.2]][0.3]
std:[[0.04, 0.04, 0.01]][[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][0.01]
MSE:[[0.29, 0.28, 0.3]][[0.3, 0.3, 0.3]][[0.2, 0.2, 0.2]][0.3]
MSE(-DR):[[0.0, -0.01, 0.01]][[0.01, 0.01, 0.01]][[-0.09, -0.09, -0.09]][0.01]
***** BETTER THAN [IS, DR_NO_MARL] *****

0_threshold = 15

MC-based mean and std of average reward:[1.1758e+01 4.0000e-03]

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.26, 0.27, 0.22]][[0.38, 0.38, 0.37]][[11.76, 11.76, 11.76]][[0.23, 0.52]]
std:[[0.01, 0.0, 0.0]][[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][[0.0, 0.01]]
MSE:[[0.26, 0.27, 0.22]][[0.38, 0.38, 0.37]][[11.76, 11.76, 11.76]][[0.23, 0.52]]
MSE(-DR):[[0.0, 0.01, -0.04]][[0.12, 0.12, 0.11]][[11.5, 11.5, 11.5]][[-0.03, 0.26]]
better than DR_NO_MARL

MC-based ATE = 0.04

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]
bias:[[0.37, 0.37, 0.32]][[0.53, 0.53, 0.52]][[0.04, 0.04, 0.04]][0.33]
std:[[0.01, 0.01, 0.01]][[0.0, 0.0, 0.01]][[0.0, 0.0, 0.0]][0.0]
MSE:[[0.37, 0.37, 0.32]][[0.53, 0.53, 0.52]][[0.04, 0.04, 0.04]][0.33]
MSE(-DR):[[0.0, 0.0, -0.05]][[0.16, 0.16, 0.15]][[-0.33, -0.33, -0.33]][-0.04]
better than DR_NO_MARL

time spent until now: 2.3 mins

[pattern_seed, T, sd_R] = [0, 672, 2]

max(u_0) = 27.327727595549877

0_threshold = 12

means of Order:

22.323 12.937 16.305 27.014 23.267

7.457 16.12 10.376 10.577 12.991

11.677 19.721 14.946 11.573 13.165

12.597 20.038 10.155 12.494 7.833

3.97 14.317 15.577 8.192 27.328

target policy:

1 1 1 1 1

0 1 0 0 1

0 1 1 0 1

1 1 0 1 0

0 1 1 0 1

number of reward locations: 16

0_threshold = 9

target policy:

1 1 1 1 1

0 1 1 1 1

1 1 1 1 1

1 1 1 1 0

0 1 1 0 1

number of reward locations: 21

0_threshold = 15

target policy:

1 0 1 1 1

0 1 0 0 0

0 1 0 0 0

0 1 0 0 0

0 0 1 0 1

number of reward locations: 9

1 2 3 1 2 3

0_threshold = 12

MC-based mean and std of average reward:[11.717 0.015]

Value of Behaviour policy:11.244

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]

bias:[[0.16, 0.15, 0.14]][[0.16, 0.16, 0.16]][[11.72, 11.72, 11.72]][[0.13, 0.47]]

std:[[0.01, 0.0, 0.01]][[0.03, 0.02, 0.02]][[0.0, 0.0, 0.0]][[0.01, 0.01]]

MSE:[[0.16, 0.15, 0.14]][[0.16, 0.16, 0.16]][[11.72, 11.72, 11.72]][[0.13, 0.47]]

MSE(-DR):[[0.0, -0.01, -0.02]][[0.0, 0.0, 0.0]][[11.56, 11.56, 11.56]][[-0.03, 0.31]]

better than DR_NO_MARL

=====

0_threshold = 9

MC-based mean and std of average reward:[11.523 0.016]

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]

bias:[[0.45, 0.45, 0.43]][[0.46, 0.45, 0.46]][[11.52, 11.52, 11.52]][[0.43, 0.28]]

std:[[0.02, 0.02, 0.01]][[0.02, 0.02, 0.02]][[0.0, 0.0, 0.0]][[0.01, 0.01]]

MSE:[[0.45, 0.45, 0.43]][[0.46, 0.45, 0.46]][[11.52, 11.52, 11.52]][[0.43, 0.28]]

MSE(-DR):[[0.0, 0.0, -0.02]][[0.01, 0.0, 0.01]][[11.07, 11.07, 11.07]][[-0.02, -0.17]]

better than DR_NO_MARL

MC-based ATE = -0.19

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]

bias:[[0.3, 0.29, 0.29]][[0.3, 0.29, 0.3]][[0.19, 0.19, 0.19]][0.29]

std:[[0.01, 0.02, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][0.0]

MSE:[[0.3, 0.29, 0.29]][[0.3, 0.29, 0.3]][[0.19, 0.19, 0.19]][0.29]

MSE(-DR):[[0.0, -0.01, -0.01]][[0.0, -0.01, 0.0]][[-0.11, -0.11, -0.11]][-0.01]

=====

0_threshold = 15

MC-based mean and std of average reward:[11.758 0.015]

[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]

bias:[[0.18, 0.19, 0.18]][[0.36, 0.36, 0.36]][[11.76, 11.76, 11.76]][[0.19, 0.51]]

std:[[0.02, 0.01, 0.01]][[0.03, 0.02, 0.02]][[0.0, 0.0, 0.0]][[0.0, 0.01]]

MSE:[[0.18, 0.19, 0.18]][[0.36, 0.36, 0.36]][[11.76, 11.76, 11.76]][[0.19, 0.51]]

```

MSE(-DR):[[0.0, 0.01, 0.0]][[0.18, 0.18, 0.18]][[11.58, 11.58, 11.58]][[0.01, 0.33]]
***** BETTER THAN [QV, IS, DR_NO_MARL] *****
MC-based ATE = 0.04
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]
bias:[[0.33, 0.34, 0.32]][[0.52, 0.52, 0.52]][[0.04, 0.04, 0.04]][0.33]
std:[[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][0.01]
MSE:[[0.33, 0.34, 0.32]][[0.52, 0.52, 0.52]][[0.04, 0.04, 0.04]][0.33]
MSE(-DR):[[0.0, 0.01, -0.01]][[0.19, 0.19, 0.19]][[-0.29, -0.29, -0.29]][0.0]
better than DR_NO_MARL
=====
time spent until now: 4.7 mins

```

```

-----
[pattern_seed, T, sd_R] = [1, 672, 0]

```

```

max(u_0) = 22.15193176791189
0_threshold = 12
means of Order:

```

```

21.11 8.63 8.924 7.177 15.583

4.39 22.152 8.13 12.524 9.977

19.783 4.835 9.689 9.453 17.349

7.1 10.289 7.759 11.211 13.917

7.098 17.425 15.81 13.477 15.805

```

target policy:

```

1 0 0 0 1

0 1 0 1 0

1 0 0 0 1

0 0 0 0 1

0 1 1 1 1

```

```

number of reward locations: 11
0_threshold = 9
target policy:

```

```

1 0 0 0 1

0 1 0 1 1

1 0 1 1 1

0 1 0 1 1

0 1 1 1 1

```

```

number of reward locations: 16
0_threshold = 15
target policy:

```

```

1 0 0 0 1

0 1 0 0 0

1 0 0 0 1

0 0 0 0 0

0 1 1 0 1

```

```

number of reward locations: 8
1 2 3 1 2 3

```

```

-----
0_threshold = 12
MC-based mean and std of average reward:[9.295e+00 5.000e-03]
Value of Behaviour policy:8.886
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.08, 0.08, 0.06]][[0.16, 0.16, 0.15]][[9.3, 9.3, 9.3]][[0.07, 0.41]]

```

```

std:[[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0]]
MSE:[[0.08, 0.08, 0.06]][[0.16, 0.16, 0.15]][[9.3, 9.3, 9.3]][[0.07, 0.41]]
MSE(-DR):[[0.0, 0.0, -0.02]][[0.08, 0.08, 0.07]][[9.22, 9.22, 9.22]][[-0.01, 0.33]]
better than DR_NO_MARL
=====
0_threshold = 9
MC-based mean and std of average reward:[9.2e+00 6.0e-03]
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.19, 0.18, 0.2]][[0.16, 0.15, 0.16]][[9.2, 9.2, 9.2]][[0.19, 0.31]]
std:[[0.01, 0.01, 0.02]][[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][[0.02, 0.0]]
MSE:[[0.19, 0.18, 0.2]][[0.16, 0.15, 0.16]][[9.2, 9.2, 9.2]][[0.19, 0.31]]
MSE(-DR):[[0.0, -0.01, 0.01]][[-0.03, -0.04, -0.03]][[9.01, 9.01, 9.01]][[0.0, 0.12]]
MC-based ATE = -0.1
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]
bias:[[0.26, 0.27, 0.26]][[0.31, 0.31, 0.31]][[0.1, 0.1, 0.1]][[0.26]]
std:[[0.01, 0.01, 0.01]][[0.01, 0.01, 0.01]][[0.0, 0.0, 0.0]][[0.02]]
MSE:[[0.26, 0.27, 0.26]][[0.31, 0.31, 0.31]][[0.1, 0.1, 0.1]][[0.26]]
MSE(-DR):[[0.0, 0.01, 0.0]][[0.05, 0.05, 0.05]][[-0.16, -0.16, -0.16]][[0.0]]
***** BETTER THAN [IS, DR_NO_MARL] *****
=====
0_threshold = 15
MC-based mean and std of average reward:[9.261e+00 5.000e-03]
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR/QV/IS]_NO_MF; [DR2, V_behav]
bias:[[0.23, 0.24, 0.18]][[0.36, 0.36, 0.35]][[9.26, 9.26, 9.26]][[0.19, 0.38]]
std:[[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0]]
MSE:[[0.23, 0.24, 0.18]][[0.36, 0.36, 0.35]][[9.26, 9.26, 9.26]][[0.19, 0.38]]
MSE(-DR):[[0.0, 0.01, -0.05]][[0.13, 0.13, 0.12]][[9.03, 9.03, 9.03]][[-0.04, 0.15]]
better than DR_NO_MARL
MC-based ATE = -0.03
[DR/QV/IS]; [DR/QV/IS]_NO_MARL; [DR2]
bias:[[0.16, 0.16, 0.12]][[0.2, 0.2, 0.2]][[0.03, 0.03, 0.03]][[0.12]]
std:[[0.01, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0, 0.0, 0.0]][[0.0]]
MSE:[[0.16, 0.16, 0.12]][[0.2, 0.2, 0.2]][[0.03, 0.03, 0.03]][[0.12]]
MSE(-DR):[[0.0, 0.0, -0.04]][[0.04, 0.04, 0.04]][[-0.13, -0.13, -0.13]][[-0.04]]
better than DR_NO_MARL
=====
time spent until now: 7.0 mins

```

```

[pattern_seed, T, sd_R] = [1, 672, 2]

```

```

max(u_0) = 22.15193176791189

```

```

0_threshold = 12

```

```

means of Order:

```

```

21.11 8.63 8.924 7.177 15.583

```

```

4.39 22.152 8.13 12.524 9.977

```

```

19.783 4.835 9.689 9.453 17.349

```

```

7.1 10.289 7.759 11.211 13.917

```

```

7.098 17.425 15.81 13.477 15.805

```

```

target policy:

```

```

1 0 0 0 1

```

```

0 1 0 1 0

```

```

1 0 0 0 1

```

```

0 0 0 0 1

```

```

0 1 1 1 1

```

```

number of reward locations: 11

```

```

0_threshold = 9

```

```

target policy:

```

```

1 0 0 0 1

```

```

0 1 0 1 1

```

```

1 0 1 1 1

```

0 1 0 1 1

0 1 1 1 1

number of reward locations: 16

0_threshold = 15

target policy:

1 0 0 0 1

0 1 0 0 0

1 0 0 0 1

0 0 0 0 0

0 1 1 0 1

number of reward locations: 8

1 2