

Last login: Wed Apr 1 14:58:34 on ttys000
Run-Mac:~ mac\$ cd ~/.ssh
Run-Mac:~.ssh mac\$ ssh -i "Runzhe_Song_0110.pem" ubuntu@ec2-3-227-211-214.compute-1.amazonaws.com
Warning: Permanently added the ED25519 host key for IP address '3.227.211.214' to the list of known hosts.
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1060-aws x86_64)

* Documentation: <https://help.ubuntu.com>
* Management: <https://landscape.canonical.com>
* Support: <https://ubuntu.com/advantage>

System information as of Wed Apr 1 20:30:39 UTC 2020

System load: 1.17 Processes: 399
Usage of /: 55.4% of 15.45GB Users logged in: 0
Memory usage: 0% IP address for ens5: 172.31.79.27
Swap usage: 0%

* Kubernetes 1.18 GA is now available! See <https://microk8s.io> for docs or
install it with:

sudo snap install microk8s --channel=1.18 --classic

* Multipass 1.1 adds proxy support for developers behind enterprise
firewalls. Rapid prototyping for cloud operations just got easier.

<https://multipass.run/>

* Canonical Livepatch is available for installation.
- Reduce system reboots and improve kernel security. Activate at:
<https://ubuntu.com/livepatch>

93 packages can be updated.
43 updates are security updates.

Last login: Thu Mar 5 21:23:34 2020 from 107.13.161.147
ubuntu@ip-172-31-79-27:~\$ export openblas_num_threads=1; export OMP_NUM_THREADS=1; python EC2.py
16:32, 04/01; num of cores:36

Basic setting:[T, sd_0, sd_D, sd_R, sd_u_0, w_0, w_A, [M_in_R, mean_reversion, pois0, simple, u_0_u_D]] = [672, 10, 10, 10, 0.3, 0.5, 1,
[True, False, True, False, 10]]

[pattern_seed, day] = [2, 7]

max(u_0) = 197.9
0_threshold = 70
means of Order:

87.8 97.8 52.4 162.7 58.1

77.3 115.7 68.5 72.4 75.7

117.4 197.9 100.7 71.1 116.9

83.2 98.9 141.5 79.5 99.8

76.4 94.9 107.4 73.9 89.9

target policy:

1 1 0 1 0

1 1 0 1 1

1 1 1 1 1

1 1 1 1 1

1 1 1 1 1

number of reward locations: 22

0_threshold = 80

target policy:

1 1 0 1 0

0 1 0 0 0

1 1 1 0 1

1 1 1 0 1

0 1 1 0 1

number of reward locations: 15

0_threshold = 90

target policy:

0 1 0 1 0

0 1 0 0 0

1 1 1 0 1

0 1 1 0 1

0 1 1 0 0

number of reward locations: 12

0_threshold = 100

target policy:

0 0 0 1 0

0 1 0 0 0

1 1 1 0 1

0 0 1 0 0

0 0 1 0 0

number of reward locations: 8

Rep 1 DONE Rep 2 DONE Rep 3 DONE Rep 4 DONE Rep 5 DONE Rep 6 DONE Rep 7 DONE Rep 8 DONE

Value of Behaviour policy:60.797

0_threshold = 70

MC for this TARGET:[69.323, 0.132]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[2.28, 2.14, 1.38]][[5.09, -69.32, -8.53]]

std:[[1.11, 1.11, 0.19]][[0.22, 0.0, 0.15]]

MSE:[[2.54, 2.41, 1.39]][[5.09, 69.32, 8.53]]

MSE(-DR):[[0.0, -0.13, -1.15]][[2.55, 66.78, 5.99]]

=====

0_threshold = 80

MC for this TARGET:[70.884, 0.141]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[0.71, 0.54, -0.96]][[1.26, -70.88, -10.09]]

std:[[0.59, 0.58, 0.42]][[0.31, 0.0, 0.15]]

MSE:[[0.92, 0.79, 1.05]][[1.3, 70.88, 10.09]]

MSE(-DR):[[0.0, -0.13, 0.13]][[0.38, 69.96, 9.17]]

=====

0_threshold = 90

MC for this TARGET:[69.371, 0.133]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[-0.15, -0.32, -1.11]][[-0.51, -69.37, -8.57]]

std:[[0.7, 0.67, 0.63]][[0.31, 0.0, 0.15]]

MSE:[[0.72, 0.74, 1.28]][[0.6, 69.37, 8.57]]

MSE(-DR):[[0.0, 0.02, 0.56]][[-0.12, 68.65, 7.85]]

=====

0_threshold = 100

MC for this TARGET:[68.94, 0.132]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[-2.7, -2.83, -3.42]][[-4.84, -68.94, -8.14]]

std:[[0.48, 0.46, 0.33]][[0.3, 0.0, 0.15]]

MSE:[[2.74, 2.87, 3.44]][[4.85, 68.94, 8.14]]

MSE(-DR):[[0.0, 0.13, 0.7]][[2.11, 66.2, 5.4]]

=====

[[2.54 2.41 1.39 5.09 69.32 8.53]

[0.92 0.79 1.05 1.3 70.88 10.09]

[0.72 0.74 1.28 0.6 69.37 8.57]

[2.74 2.87 3.44 4.85 68.94 8.14]]

time spent until now: 16.1 mins

[pattern_seed, day] = [2, 9]

max(u_0) = 197.9

0_threshold = 70

means of Order:

87.8 97.8 52.4 162.7 58.1

77.3 115.7 68.5 72.4 75.7
117.4 197.9 100.7 71.1 116.9
83.2 98.9 141.5 79.5 99.8
76.4 94.9 107.4 73.9 89.9

target policy:

1 1 0 1 0
1 1 0 1 1
1 1 1 1 1
1 1 1 1 1
1 1 1 1 1

number of reward locations: 22

0_threshold = 80

target policy:

1 1 0 1 0
0 1 0 0 0
1 1 1 0 1
1 1 1 0 1
0 1 1 0 1

number of reward locations: 15

0_threshold = 90

target policy:

0 1 0 1 0
0 1 0 0 0
1 1 1 0 1
0 1 1 0 1
0 1 1 0 0

number of reward locations: 12

0_threshold = 100

target policy:

0 0 0 1 0
0 1 0 0 0
1 1 1 0 1
0 0 1 0 0
0 0 1 0 0

number of reward locations: 8

Rep 1 DONE Rep 2 DONE Rep 3 DONE Rep 4 DONE Rep 5 DONE Rep 6 DONE Rep 7 DONE Rep 8 DONE

Value of Behaviour policy:60.741

0_threshold = 70

MC for this TARGET:[69.306, 0.104]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[2.8, 2.66, 1.61]][[5.2, -69.31, -8.56]]

std:[[0.97, 0.98, 0.49]][[0.22, 0.0, 0.1]]

MSE:[[2.96, 2.83, 1.68]][[5.2, 69.31, 8.56]]

MSE(-DR):[[0.0, -0.13, -1.28]][[2.24, 66.35, 5.6]]

=====

0_threshold = 80

MC for this TARGET:[70.872, 0.103]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[0.18, 0.01, -0.99]][[1.25, -70.87, -10.13]]

std:[[0.85, 0.83, 0.44]][[0.21, 0.0, 0.1]]

MSE:[[0.87, 0.83, 1.08]][[1.27, 70.87, 10.13]]

MSE(-DR):[[0.0, -0.04, 0.21]][[0.4, 70.0, 9.26]]

=====

```

0_threshold = 90
MC for this TARGET:[69.356, 0.108]
[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]
bias:[[-0.25, -0.41, -1.11]][[-0.5, -69.36, -8.61]]
std:[[0.26, 0.26, 0.28]][[0.26, 0.0, 0.1]]
MSE:[[0.36, 0.49, 1.14]][[0.56, 69.36, 8.61]]
MSE(-DR):[[0.0, 0.13, 0.78]][[0.2, 69.0, 8.25]]
***
=====

```

```

0_threshold = 100
MC for this TARGET:[68.927, 0.115]
[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]
bias:[[-2.9, -3.01, -3.26]][[-4.82, -68.93, -8.19]]
std:[[0.4, 0.4, 0.2]][[0.27, 0.0, 0.1]]
MSE:[[2.93, 3.04, 3.27]][[4.83, 68.93, 8.19]]
MSE(-DR):[[0.0, 0.11, 0.34]][[1.9, 66.0, 5.26]]
***
=====

```

```

[[ 2.54  2.41  1.39  5.09 69.32  8.53]
 [ 0.92  0.79  1.05  1.3  70.88 10.09]
 [ 0.72  0.74  1.28  0.6  69.37  8.57]
 [ 2.74  2.87  3.44  4.85 68.94  8.14]]

```

```

[[ 2.96  2.83  1.68  5.2  69.31  8.56]
 [ 0.87  0.83  1.08  1.27 70.87 10.13]
 [ 0.36  0.49  1.14  0.56 69.36  8.61]
 [ 2.93  3.04  3.27  4.83 68.93  8.19]]

```

time spent until now: 32.2 mins

```

-----
[pattern_seed, day] = [2, 11]

```

```

max(u_0) = 197.9
0_threshold = 70
means of Order:

```

```

87.8 97.8 52.4 162.7 58.1

```

```

77.3 115.7 68.5 72.4 75.7

```

```

117.4 197.9 100.7 71.1 116.9

```

```

83.2 98.9 141.5 79.5 99.8

```

```

76.4 94.9 107.4 73.9 89.9

```

target policy:

```

1 1 0 1 0

```

```

1 1 0 1 1

```

```

1 1 1 1 1

```

```

1 1 1 1 1

```

```

1 1 1 1 1

```

number of reward locations: 22

```

0_threshold = 80
target policy:

```

```

1 1 0 1 0

```

```

0 1 0 0 0

```

```

1 1 1 0 1

```

```

1 1 1 0 1

```

```

0 1 1 0 1

```

number of reward locations: 15

```

0_threshold = 90
target policy:

```

```

0 1 0 1 0

```

```

0 1 0 0 0

```

```

1 1 1 0 1

```

0 1 1 0 1

0 1 1 0 0

number of reward locations: 12

0_threshold = 100

target policy:

0 0 0 1 0

0 1 0 0 0

1 1 1 0 1

0 0 1 0 0

0 0 1 0 0

number of reward locations: 8

Rep 1 DONE Rep 2 DONE Rep 3 DONE Rep 4 DONE Rep 5 DONE Rep 6 DONE Rep 7 DONE Rep 8 DONE

Value of Behaviour policy:60.692

0_threshold = 70

MC for this TARGET:[69.325, 0.091]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[3.29, 3.19, 1.5]][[5.23, -69.32, -8.63]]

std:[[0.5, 0.51, 0.4]][[0.21, 0.0, 0.1]]

MSE:[[3.33, 3.23, 1.55]][[5.23, 69.32, 8.63]]

MSE(-DR):[[0.0, -0.1, -1.78]][[1.9, 65.99, 5.3]]

=====

0_threshold = 80

MC for this TARGET:[70.892, 0.092]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[0.68, 0.51, -0.91]][[1.45, -70.89, -10.2]]

std:[[0.58, 0.58, 0.26]][[0.22, 0.0, 0.1]]

MSE:[[0.89, 0.77, 0.95]][[1.47, 70.89, 10.2]]

MSE(-DR):[[0.0, -0.12, 0.06]][[0.58, 70.0, 9.31]]

=====

0_threshold = 90

MC for this TARGET:[69.378, 0.097]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[0.21, 0.07, -0.96]][[-0.34, -69.38, -8.69]]

std:[[0.36, 0.38, 0.24]][[0.27, 0.0, 0.1]]

MSE:[[0.42, 0.39, 0.99]][[0.43, 69.38, 8.69]]

MSE(-DR):[[0.0, -0.03, 0.57]][[0.01, 68.96, 8.27]]

=====

0_threshold = 100

MC for this TARGET:[68.934, 0.096]

[DR/QV/IS]; [DR_NO_MARL, DR_NO_MF, V_behav]

bias:[[-2.58, -2.7, -3.29]][[-4.69, -68.93, -8.24]]

std:[[0.58, 0.58, 0.34]][[0.27, 0.0, 0.1]]

MSE:[[2.64, 2.76, 3.31]][[4.7, 68.93, 8.24]]

MSE(-DR):[[0.0, 0.12, 0.67]][[2.06, 66.29, 5.6]]

=====

[[2.54 2.41 1.39 5.09 69.32 8.53]
[0.92 0.79 1.05 1.3 70.88 10.09]
[0.72 0.74 1.28 0.6 69.37 8.57]
[2.74 2.87 3.44 4.85 68.94 8.14]]

[[2.96 2.83 1.68 5.2 69.31 8.56]
[0.87 0.83 1.08 1.27 70.87 10.13]
[0.36 0.49 1.14 0.56 69.36 8.61]
[2.93 3.04 3.27 4.83 68.93 8.19]]

[[3.33 3.23 1.55 5.23 69.32 8.63]
[0.89 0.77 0.95 1.47 70.89 10.2]
[0.42 0.39 0.99 0.43 69.38 8.69]
[2.64 2.76 3.31 4.7 68.93 8.24]]

time spent until now: 48.4 mins

```

[pattern_seed, day] = [2, 14]

max(u_0) = 197.9
0_threshold = 70
means of Order:

87.8 97.8 52.4 162.7 58.1

77.3 115.7 68.5 72.4 75.7

117.4 197.9 100.7 71.1 116.9

83.2 98.9 141.5 79.5 99.8

76.4 94.9 107.4 73.9 89.9

target policy:

1 1 0 1 0

1 1 0 1 1

1 1 1 1 1

1 1 1 1 1

1 1 1 1 1

number of reward locations: 22
0_threshold = 80
target policy:

1 1 0 1 0

0 1 0 0 0

1 1 1 0 1

1 1 1 0 1

0 1 1 0 1

number of reward locations: 15
0_threshold = 90
target policy:

0 1 0 1 0

0 1 0 0 0

1 1 1 0 1

0 1 1 0 1

0 1 1 0 0

number of reward locations: 12
0_threshold = 100
target policy:

0 0 0 1 0

0 1 0 0 0

1 1 1 0 1

0 0 1 0 0

0 0 1 0 0

number of reward locations: 8
Rep 1 DONE

```